


THE COMMUNITY GUIDE



Keep the shot. Clean the sound.

SoundShredder helps you clean unwanted audio from a generation you want to keep. Drop in a video or audio file, choose what stays, audition the result, and download the mix or isolated tracks.



A familiar starting point: Remove music keeps dialogue and sound effects in your mix.

Unwanted generation audio

Use with exported Seedance or Genjutsu clips when the picture works but extra voices, music or effects distract.

A soundtrack you want to replace

Remove background music from an MP4 or MP3, then build your own score around the remaining audio.

Effects you want to keep

Isolate movement, impacts and ambience for your edit, or remove the effects layer when it is unwanted.

Persistent bubble sounds

Use Bubble FX for targeted cleanup, with aggressive multi-pass for stubborn Water bubbles.

ONE FILE IN. MORE CONTROL OUT.

Local processing. Windows and Mac launchers. CPU or NVIDIA GPU. Audio stays on your computer.

Download [SoundShredder on GitHub](#)

Made by **cyr4x** for the **Higgsfield Community**.

01 / GET STARTED

Download. Launch. Drop.


Get the latest Windows or Mac ZIP from [GitHub Releases](#). Extract the entire folder before launching. These are Python packages with setup launchers; Python is installed separately.

WINDOWS

Install **Python 3.12** with **Add Python to PATH**. Run **Start SoundShredder.bat**. Use the CPU or NVIDIA setup launcher if you want to choose the runtime explicitly.

MAC

On **macOS 12 or newer**, install **Python 3.11** using the universal2 installer. Run **Start SoundShredder.command**. If Finder refuses it, run **bash** followed by the dragged launcher path in Terminal. Read **START HERE - MAC.txt**.



Drop an MP4 directly into Bring your audio. Power your session controls CPU/GPU processing.

- 1 Export your clip first**
Save your Higgsfield generation locally. Drag one video/MP4 or audio/MP3 onto the drop zone, or choose **browse files**. No separate audio extraction is needed.
- 2 Choose the processing device**
Auto uses a compatible NVIDIA GPU when available, otherwise CPU. Mac uses CPU. The fallback checkbox retries on CPU if GPU memory fills up.
- 3 Choose a preset and process**
Select the cleanup you need, then click **Separate audio** or **Clean up bubbles**. Leave the launcher window open.


FIRST RUN

The app installs dependencies and downloads model weights: about 426 MB for layer separation and 1.2 GB for Bubble FX. Later processing can work offline. CPU runs can take longer. Mac hardware validation is still outstanding; Apple Metal acceleration is not implemented.

02 / CHOOSE WHAT STAYS

Start with the right preset.

Presets set the dialogue, music and effects levels for you. You can still change the sliders: **0% removes a layer**, **100% keeps it**, and values between them reduce it.



Example: music at 0%, dialogue and effects at 100%. The three layer cards control the exported mix.

PRESET	WHAT IT DOES	WHEN TO USE IT
Remove music	Keeps dialogue + effects	Replace a generated score while retaining scene audio.
Dialogue only	Keeps dialogue	Extract voices for editing or a dialogue-focused mix.
Effects only	Keeps sound effects	Reuse impacts, movement and ambience.
Remove effects	Keeps dialogue + music	Reduce unwanted SFX across the effects layer.
Remove dialogue	Keeps music + effects	Reduce unwanted voices or narration.
Bubble FX	Targets bubble-like sounds	Reduce a bubble sound while working from the original mix.

- 1

Listen, then refine

After separation, adjust the layer levels and click **Update mix**. Normal layer changes reuse the separated tracks, so another model run is not needed.
- 2

Check the category before removing it

Effects includes ambience, footsteps and impacts. Dialogue removal is not an individual-speaker selector. A layer preset cannot pick out every arbitrary noise by itself.

EXAMPLE: UNWANTED VOICES IN A GENJUTSU CLIP

Try **Remove dialogue** to retain the music and effects. If you want only the scene effects, try **Effects only**. Listen for any wanted breaths or vocal details that the separation also affects.

03 / TARGET A BUBBLE SOUND

Less bubble. More scene.

Choose **Bubble FX** when a generation has a distracting bubble, pop, gurgle or cartoon-like sound. This mode estimates that sound in the original mix instead of removing the entire effects layer.

The screenshot shows the 'Less bubble. More scene.' interface. At the top, it says 'TARGET A SOUND · BETA'. Below this is a circular icon with a green dot and the title 'Less bubble. More scene.' followed by a description: 'Reduce bubble-like sounds in the original mix. Similar pops or liquid effects may also change—check the removed-sound preview.' Below the description is a 'SOUND TO TARGET' section with a dropdown menu set to 'Water bubbles'. To the right is a 'REDUCTION' slider set to 85%, with a note 'Lower it if wanted effects begin to disappear.' Below the slider is a note: 'Changing the sound type requires a new separation.' Underneath is a section for 'Aggressive multi-pass' with a checkbox checked and a description: 'Catch stubborn bubbles by analyzing the cleaned audio again. Each pass applies your reduction level. Extra passes take longer and may reduce wanted effects.' To the right of this is a 'CLEANUP PASSES' section with a button that says '2 passes · Start here'. At the bottom is a 'Limit cleanup to a time range' section with a checkbox checked and input fields for 'From 7 sec' and 'To 11 sec'. A footer note states: 'Audio outside your selected range stays unchanged. Works on CPU or NVIDIA GPU; first use downloads a separate 1.2 GB model.'

Settings example: Water bubbles, 85% reduction, two passes, limited to 7-11 seconds.

1 Match the sound

Choose **Water bubbles**, **Bubbling & popping**, **Liquid gurgle** or **Cartoon boing**. Start with the closest description.

2 Limit the range when possible

If the problem happens only from 0:07 to 0:11, check **Limit cleanup to a time range**, enter **7** and **11**, then clean up. Leave it unchecked for the whole clip.

3 Audition the default, then adjust

Start around **85%** reduction. Compare the original and cleaned mix, then listen to **Removed sounds**. Lower the reduction or narrow the range if wanted effects are being caught.


PROTECT THE REST OF THE CLIP

Decoded audio outside the selected interval stays unchanged in the exported float WAV. Inside it, similar pops, liquid effects or overlapping sounds may also be reduced. Bubble FX is experimental.

04 / STUBBORN WATER BUBBLES

Give cleanup another pass.

Aggressive multi-pass is especially useful when a first cleanup leaves **persistent bubble remnants**. Every pass analyzes the audio left by the previous pass, giving the model another opportunity to find the remaining Water bubbles.



Enable Aggressive multi-pass under Water bubbles. Start with two passes before trying three or four.

PASS 1

Clean the original

PASS 2

Recheck what remains

PASS 3 / 4

Optional stronger cleanup

- 1 Use 2 passes as your starting point**
Choose **Water bubbles**, enable **Aggressive multi-pass**, and select **2 passes**. Keep the range tight when you know where the unwanted sound occurs.
- 2 Escalate only if the clip needs it**
Try **3 passes** for stronger cleanup or **4 passes** for the maximum option. The chosen reduction level applies on each pass. Extra passes take longer on both GPU and CPU.
- 3 Listen to the combined removal**
Removed sounds contains everything taken out across the passes. If it includes wanted impacts, ambience or vocal details, reduce the strength or pass count.
- 4 Rerun from the saved original**
Click **Clean up again** to use the saved original in a new session. Changing pass count, sound type, or multi-pass strength/range requires fresh processing. The earlier result remains available.

MORE PASSES MEAN STRONGER CLEANUP, NOT A GUARANTEE

The aim is to reduce stubborn bubbles while keeping the scene useful. More aggressive settings can also remove wanted effects. Compare the original, cleaned and removed tracks before choosing your final version.

Single-pass shortcut: with one pass, strength and range changes can use **Update mix**. Aggressive multi-pass is available for **Water bubbles only**.

05 / LISTEN AND DOWNLOAD

Every layer, its own track.

The **Listen to each layer** module presents four separate strips: **Dialogue**, **Music**, **Sound effects** and **Removed sounds**. Each has a waveform, player, seek control, preview volume/mute and an individual WAV download.



A completed Bubble FX session with all four listening tracks prepared. A quiet or flat track may simply contain little of that sound.

COMPARE THE SAME MOMENT

Playing a track pauses the others. Leave **Keep playback position when switching tracks** enabled for A/B comparison. Turn it off for independent positions. Preview volume does not change the exported mix.

BUBBLE SESSION? PREPARE THE LAYERS

Click **Prepare isolated tracks** if dialogue/music/effects are missing. They are separated from the **original audio before bubble cleanup**. Removed sounds comes from the cleanup; preparing layers preserves that result.

Normal layer separation prepares the listening tracks automatically. **Removed sounds** contains the portions excluded by your latest exported mix. If every layer is kept at 100%, that track can be silent.

A SIMPLE LISTENING ROUTINE

Play the original to hear the problem. Play the cleaned mix to judge the result. Then audition Removed sounds: it should mostly contain the audio you meant to take away. Download any track you want to keep working with.

06 / TAKE IT INTO YOUR EDIT

Download exactly what you need.

Use the cleaned mix as your replacement audio, or download separate tracks for more control in your editor. The files retain the source duration so you can align them to the original clip.



A completed four-pass Bubble FX example: preview the cleaned mix, download the WAV, or save the cleanup ZIP.

DOWNLOAD	WHAT YOU GET
Download cleaned WAV	The kept audio, ready to import into your video editor.
Download on a track	That isolated layer, or the audio removed by the latest export.
Normal separation ZIP	Dialogue, music, effects, cleaned mix, removed audio and report.
Bubble cleanup ZIP	Cleaned audio, combined removed sound and report. Download any separately prepared layers individually.

BACK INTO YOUR EDIT

Import the cleaned WAV into Premiere Pro, DaVinci Resolve or your editor. Align it to the original clip start and mute the original clip audio. Check sync and transitions. Video inputs produce audio outputs; SoundShredder does not export a replacement MP4.


- 1

Build your final soundtrack
Layer in your own score, foley or replacement dialogue. Separate downloads let you adjust each element on its own track.
- 2

Keep the processing report
The report records how the session was processed. Save it with your WAVs when comparing different strengths, ranges or pass counts.

07 / THE LEFT PANEL

Start fresh. Keep your work.



Session controls in the desktop sidebar. Example filename simplified for readability.

1. Start a new session

Click **Audio separator** at the top of the left panel. It clears the current workspace so you can drop in another file. It does **not** delete your completed saved sessions.

2. Reopen an earlier result

Click a filename under **Recent sessions** to restore its saved settings, previews and downloads. Use the refresh icon beside the heading to refresh the list.

3. Understand the x

The **x** beside a filename **deletes that local session** after confirmation. It is not a harmless close-tab button. The uploaded copy and generated files are removed; your original source file outside the app remains untouched.

4. Leave a session without deleting it

Use **Audio separator** to start fresh, or close the browser tab. Completed sessions stay on disk. Save any mix adjustments with **Update mix** first; unsaved slider changes are not a new export.

IF PROCESSING IS RUNNING

Wait for completion or click **Cancel** before starting another job or switching sessions. Closing a browser tab does not stop the local server or necessarily cancel processing.

WHEN YOU ARE FINISHED

Close the launcher/Terminal window or press **Ctrl+C** (Control+C on Mac) to stop the server. An active job is cancelled on orderly shutdown. Completed sessions remain available after restart.

KEEP DOWNLOADS BEFORE DELETING

Download important tracks to your project folder first. **Delete this session** in the results area has the same destructive purpose as the sidebar x. Session data remains in the app folder until you delete it.

The recent-session list appears in the desktop sidebar. On narrow phone-sized layouts the sidebar is hidden; use a wider desktop browser for those controls.

08 / QUICK REFERENCE

A few things worth knowing.

SUPPORTED INPUTS

MP4, MOV and other supported video containers; MP3, WAV, FLAC, M4A, AAC, OGG, Opus, AIFF and WMA. Use one file at a time. Limits: **500 MB, 10 minutes, mono or stereo.**

WHAT STAYS LOCAL

Audio processing and session files stay on your computer. First-time setup needs the internet for dependencies and models. Update checks contact GitHub only when clicked.

THE CLEANUP TOOK TOO MUCH

Lower the layer reduction or bubble strength, shorten the selected range, or use fewer passes. Listen to Removed sounds for clues. AI separation can affect breaths, singing and overlapping effects.

A RUN SEEMS SLOW


Check the progress message; first use may be downloading weights. CPU can take much longer. Use NVIDIA GPU when available, and keep the launcher window open.

PREVIEWS DO NOT MATCH NEW SETTINGS

Apply changes with **Update mix**, or **Clean up again** when fresh inference is required. The app labels old previews until you update them.

SAVING AND REUSING AUDIO

Layer exports are 24-bit WAV; Bubble FX uses 32-bit float WAV. Preserve your original source and listen in the final edit. Use source material you have permission to work with.



Update controls detail in the compact layout. On desktop, this card sits at the bottom of the left sidebar; on smaller screens, at the bottom of the workspace.

Updating is manual: check for a new release, download the matching ZIP, close the app and extract into a new folder. Copy your old **data** folder into the new one while the app is closed to retain sessions. Run the new launcher and keep the old folder until everything is checked.

GET SOUNDSHREDDER

Latest Windows and Mac downloads

github.com/xD4O/SoundShredder

Made by **cyr4x**. Made for the **Higgsfield Community**.

Follow **cyr4x**: [X](#) / [Higgsfield](#) / [Instagram](#) / [YouTube](#)

09 / WINDOWS PC INSTALLATION

Install on your Windows PC.

The public release is a Python source package with setup launchers. Install Python separately. Standalone installer previews are being tested separately.

1 Download and extract

Visit [GitHub Releases](#) and download **SoundShredder.zip**. Right-click it and choose **Extract All**. Keep the complete extracted folder in a writable location such as Documents. Do not run the launcher inside the ZIP.

2 Install Python

Install **Python 3.12** from [python.org](#). Enable **Add Python to PATH** in the installer. Reopen the launcher after installation.

3 Choose your setup

Start SoundShredder.bat installs dependencies on first launch. To choose explicitly, run **Setup CPU.bat** or **Setup NVIDIA GPU.bat** first, then the Start launcher. NVIDIA mode needs a compatible NVIDIA GPU and current driver; a separate CUDA Toolkit is not required.

4 Allow first-launch downloads

Keep an internet connection during setup. The provided NVIDIA PyTorch runtime is about **3.2 GB**. Models download when first needed: about **426 MB** for track separation and **1.2 GB** for Bubble FX. Later runs reuse the cache.

5 Open and use the app

The launcher opens your browser, normally at **http://127.0.0.1:7860**. Keep its window open. Drop in a video or audio file, choose a preset and process. CPU works without an NVIDIA GPU but takes longer.

6 Close, reopen or troubleshoot

Finish processing before closing the launcher. Reopen with **Start SoundShredder.bat**. If Python is not found, check its PATH option. If setup fails, retain the error text and try **Setup CPU.bat**, then launch again. Run one instance per session folder.

UPDATING SAFELY

Close the old app, extract the new release into a fresh folder, and copy the old **data** folder into it to retain sessions. Run the new launcher. Keep the old folder until the new installation is checked.

10 / VIDEO PREVIEW

Watch the shot. Hear each layer.

Video uploads include **Your footage, in sync**. Judge unwanted sounds against movement, dialogue and effects in the picture.

- 1 Show or hide the monitor**
Enable **Show video preview**. Your browser remembers the preference. Audio-only uploads hide the panel. Saved video sessions reopen the footage without another upload.
- 2 Choose what you hear**
Under **Listen to**, choose **Original**, **Cleaned mix**, **Dialogue**, **Music**, **Sound effects**, or **Removed sounds**. Choices become available as their tracks are prepared.
- 3 Play, pause and seek**
Use the monitor controls and timeline. Playing an existing audio player also makes the footage follow that track. Video stays muted to avoid doubling the soundtrack. Hiding it leaves audio playback available.
- 4 Check what was removed**
Compare the cleaned mix with Removed sounds at the same moment. For Bubble FX, prepare isolated tracks separately to also hear dialogue, music and effects.

PREVIEW, NOT VIDEO EXPORT

Downloads remain WAV audio and ZIP bundles. Replace the soundtrack in your video editor. Browser codec support varies; try an H.264 MP4 if preview fails. Audio cleanup can still work. Check alignment in your editor for containers with audio/video start offsets.

LOCAL PREVIEWS: EASIER INSTALLATION

Windows and separate Apple Silicon/Intel Mac standalone previews include Python and a matching setup wizard. They are unsigned; Mac hardware testing is outstanding. The public v1.0.3 download remains the source ZIP.

NEXT: SIGNING AND UPDATES

Previews add guided setup, private storage and copyable diagnostics. Signing, Mac hardware validation and automatic app updates remain future work. No public installer release date is promised.

See the [project roadmap](#). Mac installation remains in Get Started and the included **START HERE - MAC.txt** file.

11 / MAC TROUBLESHOOTING

The page opens. Processing hits an SSL error.

For the **Mac source ZIP**, including Python **3.11.9**. An error containing **CERTIFICATE_VERIFY_FAILED** or **unable to get local issuer certificate** can mean Python is missing trusted HTTPS certificates.

WHY LOCALHOST CAN STILL OPEN

The interface runs on your computer. First processing can still download an audio model from **Zenodo** (Bandit) or **Hugging Face** (Bubble FX). Those HTTPS connections need trusted certificates.

1 Stop the app fully

Press **Control+C** in SoundShredder's Terminal window. Closing only the browser tab leaves the server running.

2 Complete Python's certificate setup

For Python **3.11.9**, open **Applications > Python 3.11 > Install Certificates.command**. Or run in Terminal:
open "/Applications/Python 3.11/Install Certificates.command"
The folder is **3.11**, not 3.11.9. Wait for completion.

3 Restart and retry

Reopen **Start SoundShredder.command**, then retry processing your video or audio file. Existing sessions and the installed engine are retained.

ALTERNATIVE: REPAIR AN OLDER SOURCE INSTALL

Use **SoundShredder-Mac-Certificate-Fix.zip** from the accompanying **support/mac** folder, or the download in this HTML guide. Extract it and copy only **Start SoundShredder - Certificate Fix.command** beside the existing **app.py**. Stop the old app, then open the new launcher. Keep the existing **.venv** and **data** folders.

The repair uses Python's packaged CA roots only when its default store is empty; **SSL verification stays enabled**. If it still fails, retain the full Terminal error, Python version and processing mode. A VPN, proxy or company network may require an administrator-provided trusted certificate.

Standalone **.app** previews have their own Python and CA bundle: use **Setup details** for diagnostics; this source repair launcher does not go inside the .app. [Official Python Mac setup guide](#).