


THE COMMUNITY GUIDE



Keep the shot. Clean the sound.

SoundShredder helps you clean unwanted audio from a generation you want to keep. Drop in a video or audio file, choose what stays, audition the result, and download the mix or isolated tracks.



A familiar starting point: Remove music keeps dialogue and sound effects in your mix.

Unwanted generation audio

Use with exported Seedance or Genjutsu clips when the picture works but extra voices, music or effects distract.

A soundtrack you want to replace

Remove background music from an MP4 or MP3, then build your own score around the remaining audio.

Effects you want to keep

Isolate movement, impacts and ambience for your edit, or remove the effects layer when it is unwanted.

Persistent bubble sounds

Use Bubble FX for targeted cleanup, with aggressive multi-pass for stubborn Water bubbles.

ONE FILE IN. MORE CONTROL OUT.

Local processing in the Electron desktop app. Python is included. Windows supports CPU or NVIDIA; Mac uses CPU. Audio stays on your computer.

Windows Electron download / Mac Electron downloads

Made by **cyr4x** for the **Higgsfield Community**.

01 / GET STARTED

Download. Launch. Drop.


Start with the **v1.2.0 Electron preview**: the familiar interface in its own desktop window. Python is bundled; the engine downloads automatically. No separate Python installation or terminal setup is required.

WINDOWS

Download the [Windows EXE](#). Use **Windows 10/11 x64**. Install, then open SoundShredder from Start. Choose CPU or NVIDIA in setup.

MAC

Download the [matching Mac DMG](#). Use **macOS 13+**: arm64 for Apple Silicon, x64 for Intel. Drag the app to Applications, eject the DMG and open it. Mac uses CPU.



Drop an MP4 into Bring your audio. The shared workspace keeps the same controls in Electron.

1 Export your clip first

Save your Higgsfield generation locally, then drop one video/MP4 or audio/MP3 onto the app. No separate audio extraction is needed.

2 Choose the processing device

Auto uses a compatible NVIDIA GPU when available, otherwise CPU. Mac uses CPU. The fallback checkbox retries on CPU if GPU memory fills up.

3 Choose a preset and process

Select the cleanup you need, then click **Separate audio** or **Clean up bubbles**. Keep the app open while it works.

FIRST RUN / PREVIEW STATUS


Engine and model downloads need internet: about 426 MB for layer separation and 1.2 GB for Bubble FX. Later runs reuse the cache. These are unsigned previews; Mac builds are not notarized. CPU can be slow; Metal/MPS is not enabled.

Using a source ZIP instead? It requires your own Python. Follow the [Windows source guide](#) or [Mac source guide](#).

02 / CHOOSE WHAT STAYS

Start with the right preset.

Presets set the dialogue, music and effects levels for you. You can still change the sliders: **0% removes a layer**, **100% keeps it**, and values between them reduce it.



Example: music at 0%, dialogue and effects at 100%. The three layer cards control the exported mix.

PRESET	WHAT IT DOES	WHEN TO USE IT
Remove music	Keeps dialogue + effects	Replace a generated score while retaining scene audio.
Dialogue only	Keeps dialogue	Extract voices for editing or a dialogue-focused mix.
Effects only	Keeps sound effects	Reuse impacts, movement and ambience.
Remove effects	Keeps dialogue + music	Reduce unwanted SFX across the effects layer.
Remove dialogue	Keeps music + effects	Reduce unwanted voices or narration.
Bubble FX	Targets bubble-like sounds	Reduce a bubble sound while working from the original mix.

- 1

Listen, then refine

After separation, adjust the layer levels and click **Update mix**. Normal layer changes reuse the separated tracks, so another model run is not needed.
- 2

Check the category before removing it

Effects includes ambience, footsteps and impacts. Dialogue removal is not an individual-speaker selector. A layer preset cannot pick out every arbitrary noise by itself.

EXAMPLE: UNWANTED VOICES IN A GENJUTSU CLIP

Try **Remove dialogue** to retain the music and effects. If you want only the scene effects, try **Effects only**. Listen for any wanted breaths or vocal details that the separation also affects.

03 / TARGET A BUBBLE SOUND

Less bubble. More scene.

Choose **Bubble FX** when a generation has a distracting bubble, pop, gurgle or cartoon-like sound. This mode estimates that sound in the original mix instead of removing the entire effects layer.

The screenshot shows the 'Less bubble. More scene.' interface. At the top, it says 'TARGET A SOUND · BETA'. Below this is a circular icon with a green dot and the title 'Less bubble. More scene.' followed by a description: 'Reduce bubble-like sounds in the original mix. Similar pops or liquid effects may also change—check the removed-sound preview.' Below the description is a 'SOUND TO TARGET' section with a dropdown menu set to 'Water bubbles'. To the right is a 'REDUCTION' slider set to 85%, with a note 'Lower it if wanted effects begin to disappear.' Below the slider is a note: 'Changing the sound type requires a new separation.' Underneath is a section for 'Aggressive multi-pass' with a checkbox checked and a description: 'Catch stubborn bubbles by analyzing the cleaned audio again. Each pass applies your reduction level. Extra passes take longer and may reduce wanted effects.' To the right of this is a 'CLEANUP PASSES' section with a button that says '2 passes · Start here'. At the bottom is a 'Limit cleanup to a time range' section with a checkbox checked and input fields for 'From 7 sec' and 'To 11 sec'. A footer note states: 'Audio outside your selected range stays unchanged. Works on CPU or NVIDIA GPU; first use downloads a separate 1.2 GB model.'

Settings example: Water bubbles, 85% reduction, two passes, limited to 7-11 seconds.

1 Match the sound

Choose **Water bubbles**, **Bubbling & popping**, **Liquid gurgle** or **Cartoon boing**. Start with the closest description.

2 Limit the range when possible

If the problem happens only from 0:07 to 0:11, check **Limit cleanup to a time range**, enter **7** and **11**, then clean up. Leave it unchecked for the whole clip.

3 Audition the default, then adjust

Start around **85%** reduction. Compare the original and cleaned mix, then listen to **Removed sounds**. Lower the reduction or narrow the range if wanted effects are being caught.


PROTECT THE REST OF THE CLIP

Decoded audio outside the selected interval stays unchanged in the exported float WAV. Inside it, similar pops, liquid effects or overlapping sounds may also be reduced. Bubble FX is experimental.

04 / STUBBORN WATER BUBBLES

Give cleanup another pass.

Aggressive multi-pass is especially useful when a first cleanup leaves **persistent bubble remnants**. Every pass analyzes the audio left by the previous pass, giving the model another opportunity to find the remaining Water bubbles.



Enable Aggressive multi-pass under Water bubbles. Start with two passes before trying three or four.

PASS 1

Clean the original

PASS 2

Recheck what remains

PASS 3 / 4

Optional stronger cleanup

- 1 Use 2 passes as your starting point**
Choose **Water bubbles**, enable **Aggressive multi-pass**, and select **2 passes**. Keep the range tight when you know where the unwanted sound occurs.
- 2 Escalate only if the clip needs it**
Try **3 passes** for stronger cleanup or **4 passes** for the maximum option. The chosen reduction level applies on each pass. Extra passes take longer on both GPU and CPU.
- 3 Listen to the combined removal**
Removed sounds contains everything taken out across the passes. If it includes wanted impacts, ambience or vocal details, reduce the strength or pass count.
- 4 Rerun from the saved original**
Click **Clean up again** to use the saved original in a new session. Changing pass count, sound type, or multi-pass strength/range requires fresh processing. The earlier result remains available.

MORE PASSES MEAN STRONGER CLEANUP, NOT A GUARANTEE

The aim is to reduce stubborn bubbles while keeping the scene useful. More aggressive settings can also remove wanted effects. Compare the original, cleaned and removed tracks before choosing your final version.

Single-pass shortcut: with one pass, strength and range changes can use **Update mix**. Aggressive multi-pass is available for **Water bubbles only**.

05 / LISTEN AND DOWNLOAD

Every layer, its own track.

The **Listen to each layer** module presents four separate strips: **Dialogue**, **Music**, **Sound effects** and **Removed sounds**. Each has a waveform, player, seek control, preview volume/mute and an individual WAV download.



A completed Bubble FX session with all four listening tracks prepared. A quiet or flat track may simply contain little of that sound.

COMPARE THE SAME MOMENT

Playing a track pauses the others. Leave **Keep playback position when switching tracks** enabled for A/B comparison. Turn it off for independent positions. Preview volume does not change the exported mix.

BUBBLE SESSION? PREPARE THE LAYERS

Click **Prepare isolated tracks** if dialogue/music/effects are missing. They are separated from the **original audio before bubble cleanup**. Removed sounds comes from the cleanup; preparing layers preserves that result.

Normal layer separation prepares the listening tracks automatically. **Removed sounds** contains the portions excluded by your latest exported mix. If every layer is kept at 100%, that track can be silent.

A SIMPLE LISTENING ROUTINE

Play the original to hear the problem. Play the cleaned mix to judge the result. Then audition Removed sounds: it should mostly contain the audio you meant to take away. Download any track you want to keep working with.

06 / TAKE IT INTO YOUR EDIT

Download exactly what you need.

Download the cleaned mix or separate tracks for your editor. Electron opens a native **Save** dialog so you can choose a project folder. Exports retain the decoded audio duration; check alignment against the original footage.



A completed four-pass Bubble FX example: preview the cleaned mix, download the WAV, or save the cleanup ZIP.

DOWNLOAD	WHAT YOU GET
Download cleaned WAV	The kept audio, ready to import into your video editor.
Download on a track	That isolated layer, or the audio removed by the latest export.
Normal separation ZIP	Dialogue, music, effects, cleaned mix, removed audio and report.
Bubble cleanup ZIP	Cleaned audio, combined removed sound and report. Download any separately prepared layers individually.

BACK INTO YOUR EDIT

Import the cleaned WAV into Premiere Pro, DaVinci Resolve or your editor. Align it to the original clip start and mute the original clip audio. Check sync and transitions. Video inputs produce audio outputs; SoundShredder does not export a replacement MP4.

- 1

Build your final soundtrack


Layer in your own score, foley or replacement dialogue. Separate downloads let you adjust each element on its own track.
- 2

Keep the processing report

The report records how the session was processed. Save it with your WAVs when comparing different strengths, ranges or pass counts.

07 / THE LEFT PANEL

Start fresh. Keep your work.



Session controls in the desktop sidebar. Example filename simplified for readability.

1. Start a new session

Click **Audio separator** at the top of the left panel. It clears the current workspace so you can drop in another file. It does **not** delete your completed saved sessions.

2. Reopen an earlier result

Click a filename under **Recent sessions** to restore its saved settings, previews and downloads. Use the refresh icon beside the heading to refresh the list.

3. Understand the x

The **x** beside a filename **deletes that local session** after confirmation. It is not a harmless close-tab button. The uploaded copy and generated files are removed; your original source file outside the app remains untouched.

4. Leave a session without deleting it

Use **Audio separator** to start fresh, or quit Electron. Completed sessions stay saved. Save mix adjustments with **Update mix** first; unsaved slider changes are not a new export.

IF PROCESSING IS RUNNING

Finish or click **Cancel** before quitting. Electron blocks closing while audio processing or setup is active and explains why. Setup must finish; package-install cancellation is not available yet.

WHEN YOU ARE FINISHED

Close the Electron window or choose **SoundShredder > Quit SoundShredder**. Open it again from Start, Applications or the Dock. Completed sessions and engines remain; opening twice focuses the existing window.

KEEP DOWNLOADS BEFORE DELETING

Save important tracks to your project folder first. **Delete this session** and the sidebar x remove that session. Electron stores sessions outside its app folder; quitting or replacing the app preserves them.

Using a browser-based version? Closing its tab leaves the engine running. Quit the older standalone through its setup screen or Mac menu-bar control; stop source launchers with **Ctrl+C** / **Control+C**.

08 / QUICK REFERENCE

A few things worth knowing.

SUPPORTED INPUTS

MP4, MOV and other supported video containers; MP3, WAV, FLAC, M4A, AAC, OGG, Opus, AIFF and WMA. Use one file at a time. Limits: **500 MB, 10 minutes, mono or stereo.**

WHAT STAYS LOCAL

Audio processing and session files stay on your computer. First-time setup needs the internet for dependencies and models. Update checks contact GitHub only when clicked.

THE CLEANUP TOOK TOO MUCH

Lower the layer reduction or bubble strength, shorten the selected range, or use fewer passes. Listen to Removed sounds for clues. AI separation can affect breaths, singing and overlapping effects.

A RUN SEEMS SLOW


Read the progress message; first use may be downloading weights. CPU can take several minutes even for short clips. Use NVIDIA on compatible Windows PCs and keep the app open.

PREVIEWS DO NOT MATCH NEW SETTINGS

Apply changes with **Update mix**, or **Clean up again** when fresh inference is required. The app labels old previews until you update them.

SAVING AND REUSING AUDIO

Layer exports are 24-bit WAV; Bubble FX uses 32-bit float WAV. Preserve your original source and listen in the final edit. Use source material you have permission to work with.



The bottom-left version/GitHub card checks stable source releases. Electron previews use Help > Check for Electron updates.

Electron updates are manual: choose **Help > Check for Electron updates**, quit, then run the newer Windows installer or replace the Mac app in Applications. The separate profile preserves sessions and engines. The sidebar checker does not announce Electron prereleases.

GET SOUNDSHREDDER

Windows Electron release

Apple Silicon and Intel Mac release

GitHub project / Made by **cyr4x** for the **Higgsfield Community**.

Follow cyr4x: [X](#) / [Higgsfield](#) / [Instagram](#) / [YouTube](#)

09 / WINDOWS PC INSTALLATION

Install on your Windows PC.

Electron preview / Windows 10 or 11 x64. Python is included. Use the Windows installer; Mac downloads cannot run on Windows.

1 Get the installer

Open the [Windows Electron release](#) and download **SoundShredder-Electron-1.2.0-Windows-x64-Setup.exe**. The release includes Windows instructions and SHA-256 checksums.

2 Quit an older standalone

Use its setup screen's **Close SoundShredder** control. Keep the existing profile and sessions so Electron can reuse compatible engines. A browser tab alone does not stop it.

3 Install and open

Run the EXE, choose the application folder and finish. Open **SoundShredder** from Start or the desktop shortcut. Installation is per user; separate Python, PATH edits and administrator rights are not required.

4 Set up the engine

Choose **CPU** or **NVIDIA GPU**, then **Set up SoundShredder**. Allow at least **4 GiB free for CPU** or **14 GiB for NVIDIA**, plus the app, models and saved media. NVIDIA needs a compatible GPU and driver.

5 Use, close and return

Drop audio/video and process. Close the window or use **SoundShredder > Quit SoundShredder**. Finish setup or finish/cancel audio work first, then close again. Reopen the same shortcut to return to saved sessions.

SETTINGS, STORAGE AND UNINSTALL

SoundShredder > Setup and diagnostics opens engine controls. The default app folder is **%LOCALAPPDATA%\Programs\SoundShredder**; the separate profile is **%LOCALAPPDATA%\SoundShredder**. Windows Settings > Apps removes the app but retains engines, sessions and model caches.

This is an **unsigned preview**; Windows may show a publisher/SmartScreen prompt. [Windows troubleshooting and removal guide](#).

10 / VIDEO PREVIEW

Watch the shot. Hear each layer.

Video uploads include **Your footage, in sync**. Judge unwanted sounds against movement, dialogue and effects in the picture.

- 1 Show or hide the monitor**
Enable **Show video preview**. The app remembers the preference. Audio-only uploads hide the panel. Saved video sessions reopen the footage without another upload.
- 2 Choose what you hear**
Under **Listen to**, choose **Original**, **Cleaned mix**, **Dialogue**, **Music**, **Sound effects**, or **Removed sounds**. Choices become available as their tracks are prepared.
- 3 Play, pause and seek**
Use the monitor controls and timeline. Playing an existing audio player also makes the footage follow that track. Video stays muted to avoid doubling the soundtrack. Hiding it leaves audio playback available.
- 4 Check what was removed**
Compare the cleaned mix with Removed sounds at the same moment. For Bubble FX, prepare isolated tracks separately to also hear dialogue, music and effects.

PREVIEW, NOT VIDEO EXPORT

Downloads remain WAV audio and ZIP bundles. Replace the soundtrack in your video editor. Browser codec support varies; try an H.264 MP4 if preview fails. Audio cleanup can still work. Check alignment in your editor for containers with audio/video start offsets.

TESTED DESKTOP WORKFLOW

Packaged Windows, Apple Silicon and Intel apps completed CPU setup, real video/audio separation, preview, WAV export, duplicate launches and close/reopen with crash recovery on native CI.

WHAT REMAINS IN PREVIEW

Publisher signing/notarization, consumer Gatekeeper/SmartScreen testing, broader hardware/GPU validation and automatic app updates remain future work. Mac Bubble FX inference was not part of those CI checks.

Read the [verification record](#) / [Future updates](#). See the Windows and Mac installation pages in this guide for the correct desktop download.

11 / MAC TROUBLESHOOTING

The page opens. Processing hits an SSL error.

For the **Mac source ZIP**, including Python **3.11.9**. An error containing **CERTIFICATE_VERIFY_FAILED** or **unable to get local issuer certificate** can mean Python is missing trusted HTTPS certificates.

WHY LOCALHOST CAN STILL OPEN

The interface runs on your computer. First processing can still download an audio model from **Zenodo** (Bandit) or **Hugging Face** (Bubble FX). Those HTTPS connections need trusted certificates.

1 Stop the app fully

Press **Control+C** in SoundShredder's Terminal window. Closing only the browser tab leaves the server running.

2 Complete Python's certificate setup

For Python **3.11.9**, open **Applications > Python 3.11 > Install Certificates.command**. Or run in Terminal:
open "/Applications/Python 3.11/Install Certificates.command"
The folder is **3.11**, not 3.11.9. Wait for completion.

3 Restart and retry

Reopen **Start SoundShredder.command**, then retry processing your video or audio file. Existing sessions and the installed engine are retained.

ALTERNATIVE: REPAIR AN OLDER SOURCE INSTALL

Use **SoundShredder-Mac-Certificate-Fix.zip** from the accompanying **support/mac** folder, or the download in this HTML guide. Extract it and copy only **Start SoundShredder - Certificate Fix.command** beside the existing **app.py**. Stop the old app, then open the new launcher. Keep the existing **.venv** and **data** folders.

The repair uses Python's packaged CA roots only when its default store is empty; **SSL verification stays enabled**. If it still fails, retain the full Terminal error, Python version and processing mode. A VPN, proxy or company network may require an administrator-provided trusted certificate.

Electron or older standalone .app? These have private Python and CA bundles. The source repair does not repair that runtime. Use **SoundShredder > Setup and diagnostics** in Electron; keep certificate checks enabled. [Electron Mac troubleshooting](#) / [Official Python Mac setup guide](#).

12 / MAC DESKTOP INSTALLATION

Install on your Mac.

Electron preview / macOS 13 Ventura or newer. Python is bundled. Apple Silicon and Intel have separate downloads; both use CPU processing.

1 Choose your architecture

Check **Apple menu > About This Mac**. Choose **arm64** for M-series Apple Silicon, or **x64** for Intel. [Download the matching DMG or ZIP](#).

2 Quit and move the app

Quit any older standalone from its setup screen or menu-bar Quit control. Open the DMG, drag **SoundShredder** to **Applications**, replace the older app if prompted, then eject the disk image. For ZIP, extract and move the app to Applications.

3 Open and set up

Open SoundShredder from Applications. Select **Set up SoundShredder** to download the CPU engine. Keep internet connected and at least **3 GiB free** for setup, plus room for the app, models and sessions. No Homebrew or separate Python is needed.

4 Create, quit and reopen

Drop a file and process. Close the window or use **SoundShredder > Quit SoundShredder (Cmd+Q)**. Finish setup or finish/cancel audio work first, then close again. Reopen from Applications or the Dock; saved sessions remain.

5 Return to settings or update

Cmd+, opens Setup and diagnostics; **Cmd+1** returns to Workspace. For updates, use **Help > Check for Electron updates**, quit and replace the app. Compatible engines and sessions are reused.

UNSIGNED PREVIEW

This build has no Developer ID signing or notarization. For a trusted download, macOS may offer **System Settings > Privacy & Security > Open Anyway**. Do not disable Gatekeeper globally. If no override is offered or macOS reports damage, report the exact message; a signed release may be required.

Profile: **~/Library/Application Support/SoundShredder**. Quit and move the app to Trash to uninstall; sessions, engines and models remain. [Mac storage, SSL and removal guide](#).