




Keep the shot. Clean the sound.

Clean distracting audio in Seedance, Genjutsu and other generations, or work with existing MP4, MP3 and WAV files. Drop one file, choose what stays, compare the layers and download audio for your edit.



Akira example: a real completed session, with the cleaned mix synchronized to the footage.

REMOVE A GENERATED SCORE

Keep estimated dialogue and effects while making room for your own music.

CONTROL VOICES AND EFFECTS

Isolate dialogue, music or scene effects, or exclude an unwanted layer.

TARGET STUBBORN BUBBLES

Use Water bubbles and optional aggressive multi-pass for persistent remnants.

SEE AND HEAR THE RESULT

Toggle video preview, audition four tracks and save WAVs independently.

[Download Windows EXE](#) / [Download Apple Silicon DMG](#) / [Download Intel Mac DMG](#)

Made by cyr4x. Made for the Higgsfield Community.

Version 1.3.0 / Documentation refreshed September 26, 2026.

Community-created; not an official Higgsfield product.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Install on Windows.

Windows 10/11 x64 / Electron desktop app. Python is bundled. Engine and model downloads happen inside the app on first use.

WINDOWS SIGNING IS PENDING

The 1.3.0 EXE is unsigned. Edge may show an uncommon-download warning; Windows may show SmartScreen or an unknown publisher. Use the official xD4O/SoundShredder release and compare checksums. See the next page before running the installer.

1 Download the EXE

Download [Windows EXE](#) from the official release. Read the [Windows warning page](#) and verify the downloaded file.

2 Quit old copies, then install

Quit an older app or use Close SoundShredder on its standalone setup page. Run the EXE, choose the app installation folder, and finish. No separate Python, PATH edits or administrator rights are needed.

3 Open and choose storage

Open SoundShredder from Start or its desktop shortcut. Use **Choose folder...** for engines, models and sessions, or keep the current location. This storage choice is separate from the app installation folder.

4 Set up the engine

Select CPU or a compatible NVIDIA GPU, then **Set up SoundShredder**. Keep internet connected. Allow at least **4 GiB free for CPU** or **14 GiB for NVIDIA**, plus the app, models and media. NVIDIA needs a compatible driver; no separate CUDA Toolkit is needed.

5 Use it, quit and reopen

The workspace opens automatically. Finish or cancel processing before quitting. Reopen the same shortcut next time; compatible engines and saved sessions are reused.

Uninstall: quit, then use **Uninstall SoundShredder** in Start or Windows Settings > Apps. Sessions and engines are retained. [Full Windows guide and removal instructions](#).

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Understand the warning.

Our official download source is github.com/xD4O/SoundShredder/releases/latest. The Windows installer is still unsigned. The signed Mac release does not sign the Windows EXE.

COMMON MESSAGE

Edge: **isn't commonly downloaded**.

Windows: **Windows protected your PC** or an unknown publisher.

These can be reputation warnings, rather than a report of a specific detected threat.

1 Confirm the exact source

Use the xD4O/SoundShredder repository and its release Assets. Avoid reuploaded installers and lookalike projects. [Open the official release](#).

2 Match the checksum

Download the EXE and SHA256SUMS.txt from the same release. Run **Get-FileHash -Algorithm SHA256 -LiteralPath** followed by the quoted path to your downloaded EXE. Compare the entire hash against its filename in the checksum file. [Windows download verification and warning guide](#) includes a copyable command.

3 Keep protections enabled

A matching hash checks the bytes against the published file; it is not a malware scan. A named threat detection or a mismatched hash needs investigation. Do not run a mismatched download or disable security protection.

4 Choose whether to continue

Only after verifying and trusting the download, Edge may offer **Keep > Show more > Keep anyway**; Windows may offer **More info > Run anyway**. Availability varies. If blocked without an override, use the source version where permitted or contact your administrator.

WHAT WE CAN VERIFY

The official installer passed functional install, uninstall, processing and reopen tests. Those checks are not a security audit, and GitHub hosting alone does not guarantee safety. Windows signing is planned; newly signed downloads can still need reputation.

[Microsoft: how SmartScreen reputation works / Windows signing roadmap](#)

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Install on your Mac.

macOS 13 or newer / Electron desktop app. Current 1.3.0 Apple Silicon and Intel builds are Developer ID signed and Apple-notarized. Mac audio processing uses CPU; Metal/MPS is not enabled.

APPLE SILICON

Download [Apple Silicon DMG](#)

For M-series chips. Check Apple menu > About This Mac.

INTEL

Download [Intel Mac DMG](#)

For an Intel processor. Use the matching architecture.

1 Quit any older SoundShredder copy

Use its Quit command or, for older browser standalones, Close SoundShredder on the setup page. Keep your saved engine and session folders.

2 Move the app into Applications

Open the matching DMG, drag SoundShredder into Applications, replace the old app if needed, then eject the disk image. Open SoundShredder from Applications.

3 Choose storage and set up

Optionally use **Choose folder...**, then **Set up SoundShredder**. Allow at least **3 GiB free**, plus the app, models and saved media. Internet is required initially. No Homebrew, separate Python or developer membership is needed to use the app.

4 Quit and return

Finish or cancel audio work before quitting. During setup, Cancel setup and quit is available. Reopen from Applications or the Dock; completed sessions remain saved.

IF FINDER REPORTS DAMAGE

The old 1.2.0 preview was unsigned and unnotarized. Download the current 1.3.0 build, check its checksum and architecture, then copy it into Applications. If the warning persists, report the exact message. Do not disable Gatekeeper or remove quarantine as an installation step.

Uninstall: quit and move SoundShredder.app to Trash. Saved data remains in the chosen folder and the user profile. [Mac storage, removal and troubleshooting](#).

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Prefer your browser?

The Python source version runs on your own computer at localhost. It uses the same workspace. Choose this route if you want source access; Electron is the easier installation for most users.

Download the [source ZIP](#). Extract the entire folder into a writable location such as Documents. Keep all files together.

WINDOWS SOURCE

Install **Python 3.12** from python.org with **Add Python to PATH**. Run **Setup CPU.bat** or **Setup NVIDIA GPU.bat**. When finished, run **Start SoundShredder.bat**.

MAC SOURCE

Use **macOS 12+** and the **Python 3.11 universal2** installer. Run **Start SoundShredder.command**. If Finder will not run the launcher, type **bash** in Terminal, drag the file into that window and press Return.

1 Wait for initial setup

The launcher downloads the CPU or selected NVIDIA engine. First processing also downloads model weights. The browser normally opens **http://127.0.0.1:7860**.

2 Keep the launcher window open

Closing the browser tab does not stop the server. Use **Ctrl+C** on Windows or **Control+C** on Mac to stop it. Reopen the same launcher next time.

3 Repair Mac source certificates if needed

For Python 3.11.9 with CERTIFICATE_VERIFY_FAILED, run **Applications > Python 3.11 > Install Certificates.command**, then restart SoundShredder. Keep SSL verification enabled. This repairs the source Python environment, not Electron's bundled runtime.

[Complete source installation instructions](#) / [Mac source SSL troubleshooting and repair launcher](#)

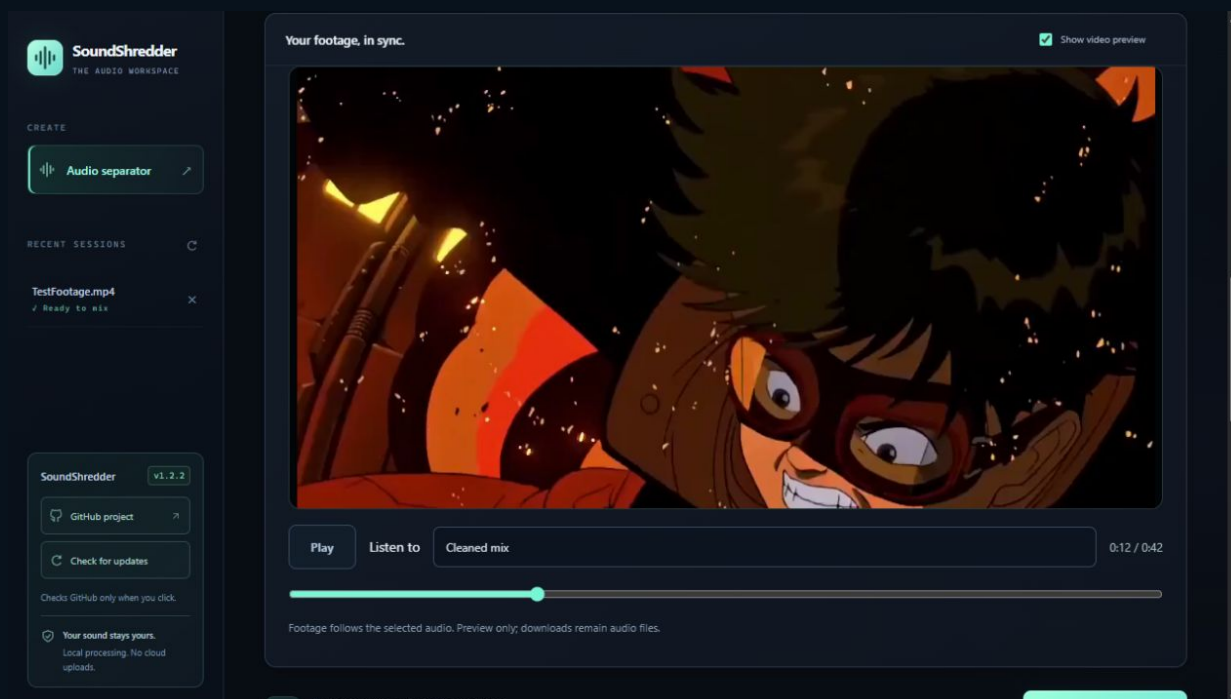
LOCAL PROCESSING

Media stays on your computer. Dependencies and model weights need internet initially; cached processing works offline. No SoundShredder account or API key is required.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Try it with Akira.

This worked example uses the supplied **TestFootage.mp4**: approximately **43 seconds**, with **48 kHz stereo** audio. The screenshots show a completed local Windows source-app run; Electron shares these controls.



Show video preview is enabled. Cleaned mix is selected at 0:12 for comparison.

1 Drop the MP4 and remove music

Choose **Remove music**: Dialogue 100%, Music 0%, Sound effects 100%. Choose Auto, CPU or NVIDIA GPU, then **Separate audio**. The example used Auto on an RTX 5090; Mac users choose CPU.

2 Compare with the footage

Enable **Show video preview**. Switch **Listen to** between Original and Cleaned mix at the same moment. The monitor uses muted video synchronized to the selected audio.

3 Refine the result

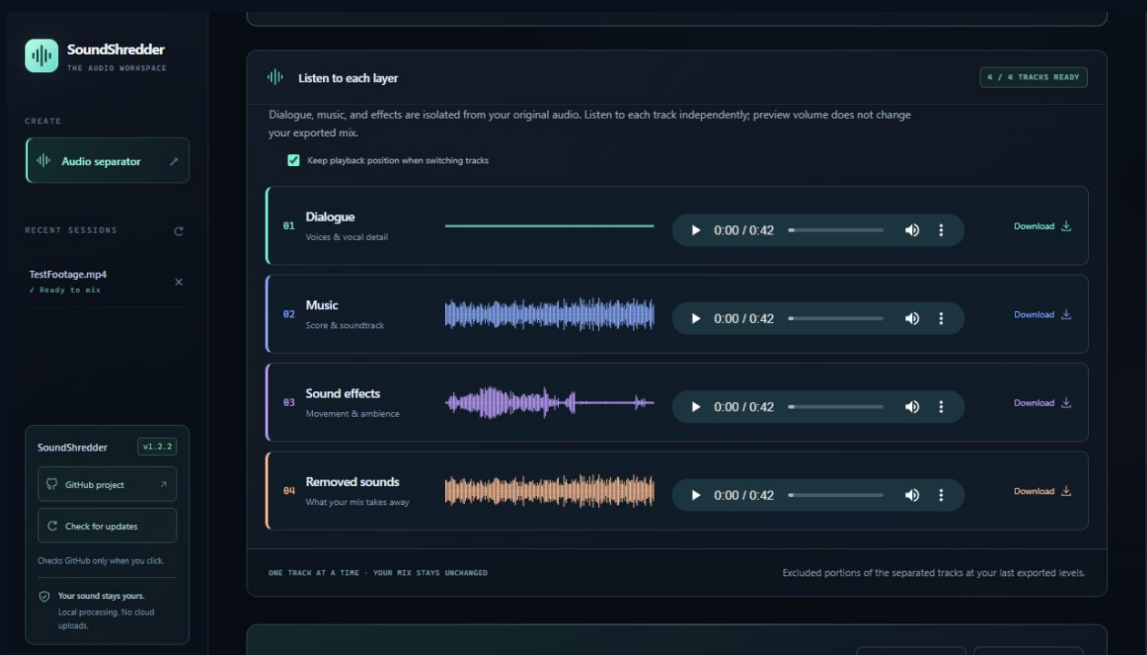
Audition wanted effects and any voices. If needed, adjust layer levels and click **Update mix**. Disable Show video preview to return to an audio-only workspace.

AN EXAMPLE, NOT A QUALITY GUARANTEE

This run demonstrates music removal. It does not establish that the Akira clip contains bubbles. Screenshots are included for the walkthrough; the full video and extracted soundtrack are not distributed. Use your own clip to follow along.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Every layer, its own track.



The Akira session: four track strips with separate playback controls and WAV downloads.

DIALOGUE + MUSIC

Dialogue is the estimated voice layer; it is nearly silent in this example. Music contains the estimated score. A quiet waveform does not necessarily indicate an error.

EFFECTS + REMOVED SOUNDS

Effects contains estimated scene audio. Removed sounds contains material excluded by your latest mix; here it corresponds to the excluded music layer.

1 Compare one track at a time

Keep **Keep playback position when switching tracks** enabled to audition the same moment. Starting a track pauses the others. Preview volume/mute does not change the export.

2 Download the layer you need

Use the Download control on Dialogue, Music, Sound effects or Removed sounds. Normal separation prepares them automatically.

3 Preparing tracks after Bubble FX

Use **Prepare isolated tracks** to hear dialogue/music/effects from the original audio before bubble cleanup. This extra separation preserves the cleaned bubble result; download its added layers individually.

LISTEN TO WHAT WAS REMOVED

If Removed sounds includes wanted impacts, ambience or vocal detail, reduce the cleanup or revise the mix. AI separation can leave bleed or reduce wanted sounds.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Choose what stays.

Layer presets set the estimated dialogue, music and effects levels. **0% excludes a layer; 100% keeps it.** Intermediate levels reduce it. Change levels after separation and click Update mix to reuse the existing stems.

PRESET	KEEPS / TARGETS	USE IT FOR
Remove music	Dialogue + effects	Replacing an unwanted generated score.
Dialogue only	Dialogue	Editing or reusing estimated voices.
Effects only	Sound effects	Working with impacts, movement and ambience.
Remove effects	Dialogue + music	Reducing an unwanted effects layer.
Remove dialogue	Music + effects	Reducing unwanted narration or voices.
Bubble FX	Selected bubble-like sound	Targeted cleanup instead of excluding all effects.

SEEDANCE / GENJUTSU

Keep a generation whose picture works but soundtrack distracts. Export the video, then drop it into SoundShredder. Choose the preset for the unwanted audio category.

OTHER VIDEO OR MP3 SOURCES

Apply the same workflow to your existing media. Dialogue removal is not an individual-speaker selector, and Effects only can include ambience.

MODEL LIMITS

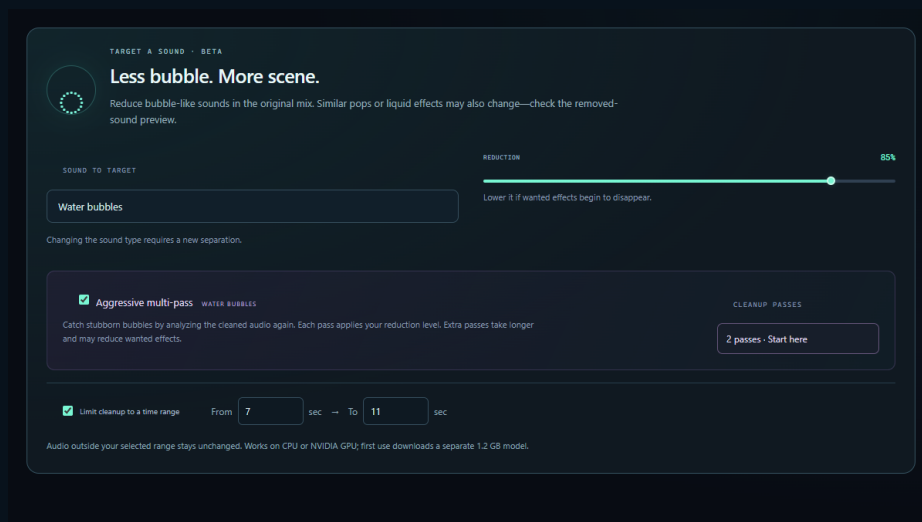
The layer model can confuse singing, breaths, rhythmic impacts and music. It cannot identify every arbitrary sound or preserve all overlapping detail. Audition the original, cleaned mix and removed material before editing.

Custom text prompts: Bubble FX currently uses four preset sound descriptions. Use Target a Sound - Guided for a free-form description, preview and saved cleanup version.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Target stubborn bubbles.

Use **Bubble FX** for a distracting bubble, pop or gurgle while trying to retain other effects. This uses AudioSep separately from the dialogue/music/effects model.



Separate bubble-settings example: 85% reduction, two passes and a 7-11 second range. Not the Akira run.

- 1 Match the unwanted sound**
Choose **Water bubbles**, **Bubbling & popping**, **Liquid gurgle** or **Cartoon boing**. Start near **85%** reduction and restrict the time range when possible.
- 2 Try aggressive multi-pass for Water bubbles**
Enable **Aggressive multi-pass** and start with **2 passes**. Each pass rechecks the previous cleaned audio. Try 3 or 4 only if stubborn remnants remain. More passes take longer and can affect similar wanted effects.
- 3 Check the combined removed sound**
Compare Original, Cleaned mix and Removed sounds. The latter combines removal across all passes. Reduce strength, narrow the interval or use fewer passes if wanted effects are caught.
- 4 Rerun when settings need fresh inference**
Changing sound type, pass count or multi-pass strength/range needs **Clean up again**, which starts from the saved original in a new session. Single-pass strength/range changes can use **Update mix**.

EXPERIMENTAL CLEANUP

The Bubble FX multi-pass switch applies to Water bubbles. Guided prompt cleanup also offers 1-4 passes for custom targets. It can reduce persistent remnants, but cannot promise bubble-only removal. Decoded samples outside the chosen interval stay unchanged in the float WAV.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Describe the distraction.

A generation has useful voices and impacts, but one repeated sound distracts. Use **Target a Sound - Guided** to describe that sound instead of removing the whole Effects layer.

1 Prepare your source

Drop the video or audio file into the separator. Switch to Target a Sound, then Prepare source. The original upload is retained.

2 Describe one sound

Try a short phrase such as **water bubbling**, **footsteps** or **a dog barking**. The prompt describes what to remove. It is not a guarantee that the model can isolate that sound in every recording.

3 Choose the interval and preview

Select where the distraction occurs, then create a short preview. Compare Before, Cleaned and Removed. If voices or wanted effects appear in Removed, narrow the interval, lower strength or rephrase.

4 Set strength and passes

Strength controls how much of the estimated target is subtracted. Passes repeat cleanup on the previous pass. Start with one; use 2-4 for stubborn remnants after checking the sound.

5 Save or continue in steps

Save a full cleanup version for export. To address another distraction afterward, explicitly use the cleaned version as the next input. Earlier versions and the original remain available.

LOCAL MODEL, NOT A PROMISE

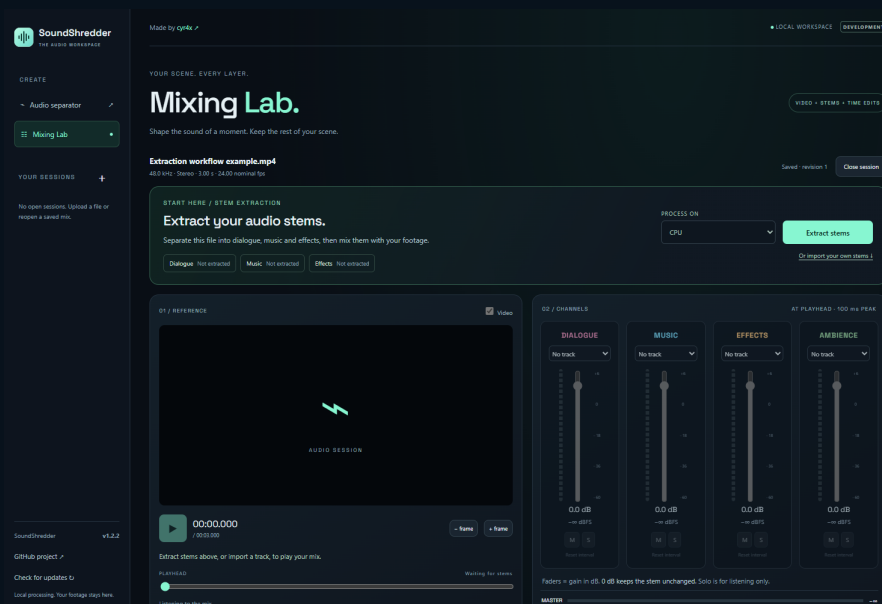
Prompt interpretation and separation can be imperfect when sounds overlap. More passes can remove wanted detail or introduce artifacts. Audition Removed as carefully as Cleaned; a successful job alone does not mean a better soundtrack.

Full targeted-cleanup guide

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Meet your Mixing Lab.

Use the Lab when you want to keep most of a scene and adjust individual moments. The reference picture and all playing stems share one playhead.



Mixing Lab demonstration: start with the Stem Extraction area above the reference monitor. This screenshot shows example UI data, not a separation-quality result.

1 Upload, then extract

Choose Mixing Lab, drop a video/audio file, choose CPU/Auto/NVIDIA where available, then click **Extract stems**. Uploading alone does not extract them. Progress, Cancel and Retry are in this area.

2 Shape one moment

Choose one second, one nominal frame or a range. Adjust Dialogue, Music or Effects at the playhead. The saved level returns outside the selected interval. 0 dB means unchanged gain; the moving meter shows signal level in dBFS.

3 Listen with the picture

Use the shared Play control, waveform or scrubber. Solo isolates a channel for listening; it does not change the exported mix. Toggle Video to hide the reference picture.

AMBIENCE

Ambience is an optional imported track, such as room tone or wind. It is not automatically separated from Effects. Recursive Split deeper is not included.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Clean a stem. Keep the scene.

If bubbles sit in Effects while dialogue and music already sound right, target the **Effects** stem in the Lab. This keeps the other selected channels out of that cleanup operation.

1 Choose a stem and interval

Select the stem and time range. Describe the unwanted sound, set strength and 1-4 passes, then preview. Compare Before, Cleaned and Removed before applying.

2 Keep versions available

Applying cleanup creates a new version for that channel. Removed audio is available to audition or download and does not automatically enter the master mix.

3 Import a replacement

Bring in an external stem or ambience track when useful. Confirm the import fitting option if its length or format differs; keep the original source as a reference.

4 Export the result

Export the automated channel WAVs, master WAV, edit report or bundle ZIP. For supported video sources, export an MP4 with the mixed soundtrack. Review the result in motion and check sync in your editor.

CLOSE A LAB SESSION

Close session and the sidebar x move it into **Closed sessions**. Audio and saved edits remain. Expand that list and select a session to reopen it.

CLOSE IS NOT CANCEL

Closing a session stops its playback but does not cancel a background job. Use Cancel extraction or Cancel processing to stop that work. Main-separator Delete is a separate destructive action.

[Complete Mixing Lab guide](#)

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Take it into your edit.

Download a ready-to-use mix or separate layers for more control. Electron uses a native Save dialog; the source web app uses your browser's download behavior.

DOWNLOAD	CONTENTS
Cleaned WAV	Audio retained by your latest settings.
Individual track	Dialogue, music, effects or removed audio.
Layer-separation ZIP	Stems, cleaned mix, removed audio and processing report.
Bubble cleanup ZIP	Cleaned audio, combined removed sound and report. Additional prepared layers are downloaded separately.

- 1

Align the WAV to the original

Import it into Premiere Pro, DaVinci Resolve or your editor. Place it at the original clip start and mute the original soundtrack. Check sync and transitions.
- 2

Keep useful isolated layers

Put music, voices and effects on their own editor tracks. Build your score or foley around the material you want to retain.
- 3

Save before deleting a session

Download the outputs to your project folder and keep the processing report. Deleting a SoundShredder session removes its local generated files.

AKIRA TIMING CHECK

The example exports preserve the decoded source: **2,052,096 frames / 48,000 Hz / 42.752 seconds / stereo**. Timing alignment was checked; these numbers do not measure audible separation quality.

The main separator exports audio. Mixing Lab can also export an MP4 with the mixed soundtrack. Layer separation exports 24-bit PCM WAV; Bubble FX exports 32-bit float WAV. Original media is preserved.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

The left panel, explained.

START A NEW SESSION

Click **Audio separator** at the top of the left panel. The workspace resets so you can import another file; completed saved sessions remain.

RETURN TO SAVED WORK

Click a filename under **Recent sessions**. The Akira example appears as **TestFootage.mp4**. Use the refresh control to update the list.

THE X DELETES A SESSION

The **x** beside a filename removes that local session after confirmation. It is not a close-tab button. **Delete this session** has the same purpose.

LEAVE WITHOUT DELETING

Start a new session or quit the app. Completed work remains. Click **Update mix** to save slider changes before leaving; unsaved changes are not a new export.

1 Finish or cancel active work

Wait for processing or click Cancel before quitting. During desktop setup, choose Cancel setup and quit if you need to stop.

2 Quit Electron and reopen normally

Close its window or use **SoundShredder > Quit SoundShredder**. Use the Start shortcut or Applications/Dock next time. Opening twice brings the existing window forward.

3 Stop source and older browser launchers correctly

Closing a browser tab leaves the server running. Stop the source launcher with Ctrl+C / Control+C; older standalone apps have a Close SoundShredder setup control.

VERSION, GITHUB AND UPDATES

The bottom-left card contains the version, GitHub link and update controls. Electron checks automatically; open the panel to download or choose Restart & update. The browser version updates manually.

SOUNDSHREDDER / COMMUNITY FIELD GUIDE

Storage, updates and help.

CHOOSE YOUR STORAGE

Use **SoundShredder > Choose storage folder....** The app remembers the choice. Files are not moved; choose your old folder again to return to its sessions. Keep external drives connected.

SETUP FEELS SLOW

Open Setup details to see package activity. Download bars can pause while packages install. Cancel and retry are available; keep enough free space and an internet connection.

YOU CHOOSE WHEN TO RESTART

Electron checks automatically. Download with progress/cancel, then Restart & update or Later. Normal Quit does not install. Active processing, uploads and saves block the restart. Coming from 1.2.2? Install 1.3.0 manually once.

FILES AND DEVICES

One file at a time: up to **500 MB / 10 minutes / mono or stereo**. Windows supports CPU or compatible NVIDIA CUDA; Mac uses CPU. AMD GPU and Apple Metal/MPS acceleration are not implemented.

Models: first use downloads about 426 MB for Bandit layer separation or 1.2 GB for Bubble FX. Cached runs can work offline. CPU processing may take longer than the clip duration.

Need help? Include the exact release filename, OS/chip, processing mode and error. Use Setup details or Open logs folder; review private paths before sharing. A named malware detection should be reported separately from an uncommon-download warning.

[Report an issue](#) / [All documentation](#) / [Verification and limits](#)

[Windows guide](#) / [Mac guide](#) / [Mac source SSL repair](#)

MADE BY CYR4X

[X](#) / [Higgsfield](#) / [Instagram](#) / [YouTube](#)

Made for the Higgsfield Community. Community-created; not an official Higgsfield product.